



Thread-Level Parallelism

Lecture 5




Thread-Level Parallelism


- ★ Threads (Revisit)
- ★ Superscalar/Multiprocessing & Threads
- ★ Fine-grained Multithreading
- ★ Simultaneous Multithreading
 - Thread-Level Parallelism to Instruction-Level Parallelism
- ★ Impact on Hardware
- ★ Software Model
 - OS Threads (POSIX Thread)
 - OpenMP
 - Apple GCD



Threads


- ★ Light-weight process
- ★ Multiple threads can exist in a process
 - Shared memory/(global) variables
 - Different PC
 - Different Stack
 - Different registers
- ★ Kernel-level Thread vs. User-level Thread
- ★ Only Kernel-level thread is visible to CPU






Superscalar/Multiprocessing and Threads

Superscalar



Multiprocessing



Thread 1




Thread 2




Thread 3

- ★ Do we have enough (benchmark) parallelism to fill superscalar and Multiprocessing?
- ★ Can we interchange thread-level parallelism and instruction-level parallelism?




Fine-grained Multithreading & Simultaneous Multithreading

**Fine grained
MT**




**Simultaneous
MT**



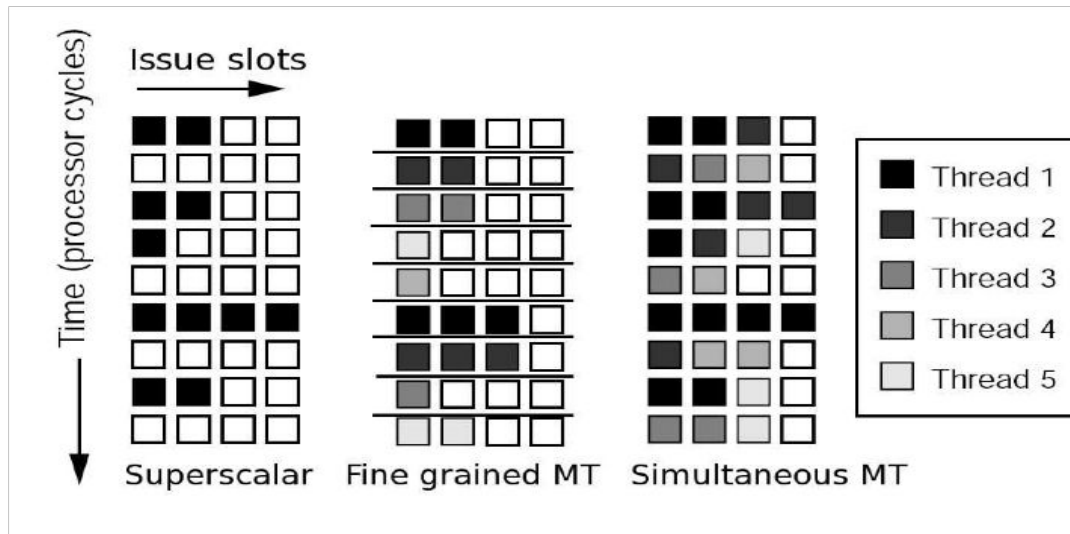
-  Thread 1
-  Thread 2
-  Thread 3


- ★ Fine-grained MT - What if we switch thread every cycle?
- ★ Simultaneous MT - How about we mixed threads in the same cycle?



Simultaneous Multithreading

Comparison






Speedup from SMT

Result from "Converting Thread-Level Parallelism to Instruction-Level Parallelism via Simultaneous Multithreading," ACM Transactions on Computer Systems, Vol. 15, No. 3, August 1997, Pages 322–354.

Table V. Throughput Comparison of MP2, MP4, and SMT, Measured in Instructions per Cycle

Configuration	Number of Threads			
	1	2	4	8
MP2	2.08	3.32	—	—
MP4	1.38	2.25	3.27	—
SMT	2.40	3.49	4.24	4.88



Impact on HW

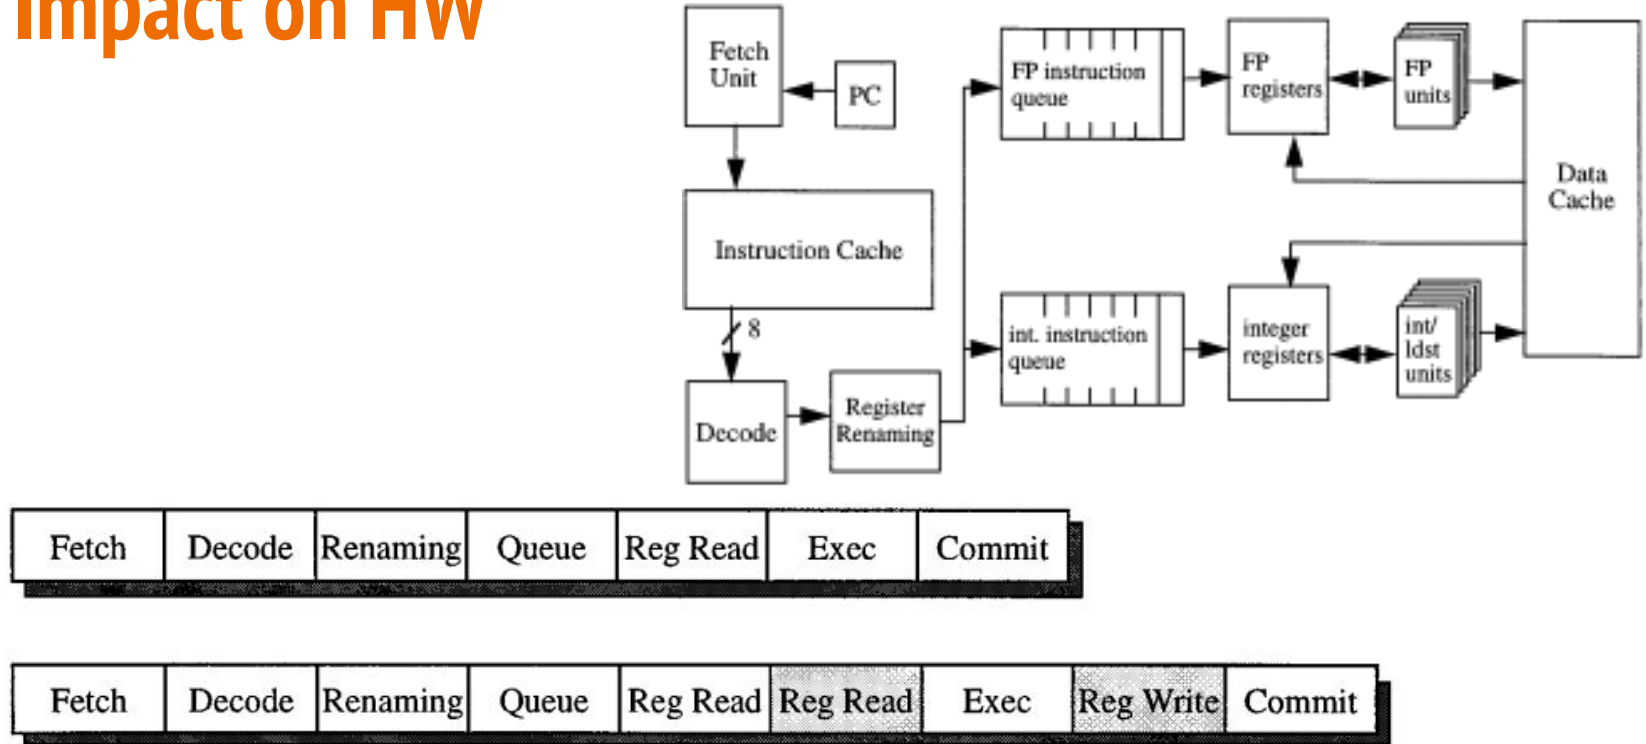

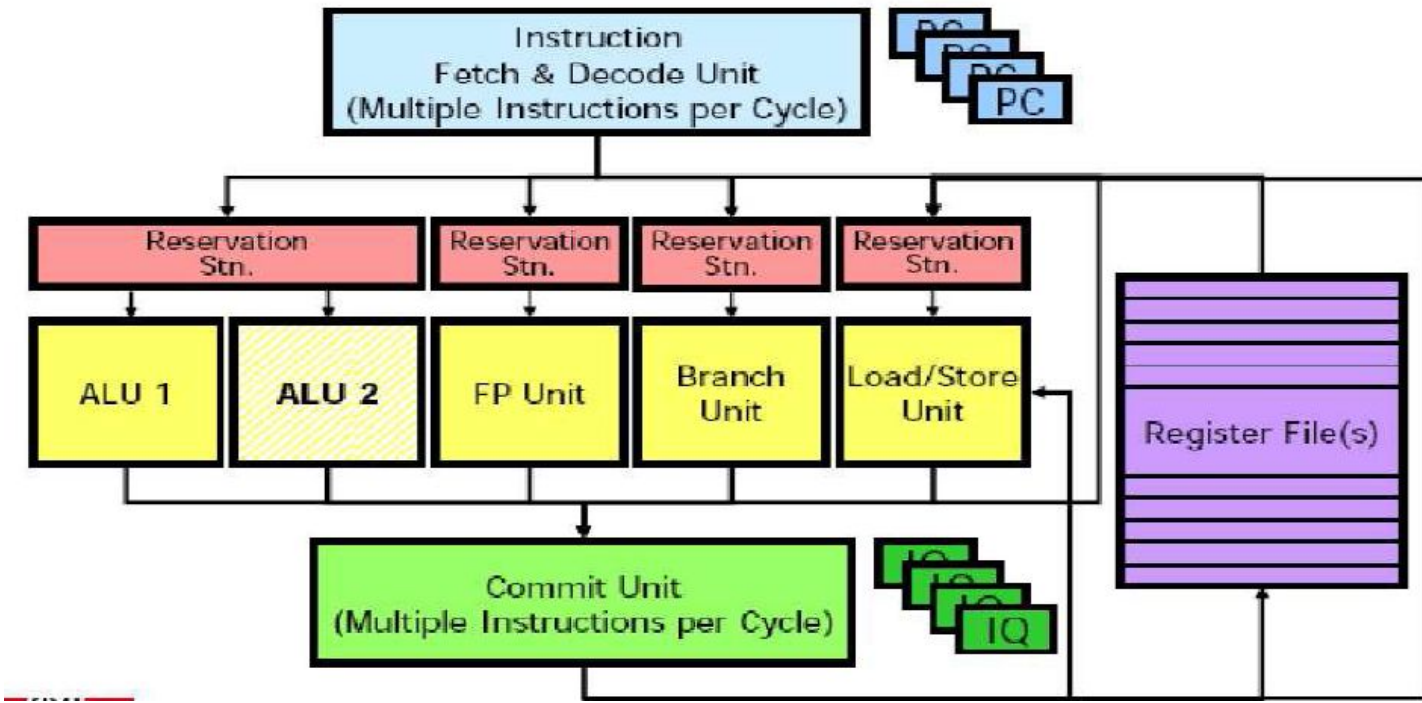




Fig. 3. Comparison of the pipelines for a conventional superscalar processor (top) and the SMT processor's modified pipeline (bottom).



Should we implement it?


- ★ (Minimal) extension to superscalar
- ★ More resources (per thread)
 - Replicated registers, PCs
 - Pipeline flushing/returns
 - Branch predictions?
- ★ Fetch Unit
- ★ Scheduling





Performance and Implementation Issues

- ★ Performance limitations
 - Issue limits - limit in functional units and availability
 - Too few renaming registers
 - Instruction queue full
- ★ Implementation Issues
 - What to fetch?
 - What to issue?
 - Cache
 - Synchronization
 - Compiler supports




Issue 1: What to fetch?

★ static

- RR
- 8 instructions from one thread
- 4 instructions from two threads
- 2 instructions from four threads


★ dynamic

- Any number of threads with minimal
 - branches?
 - misses?
 - Instructions?




Issue 2: What to issue?

- ★ Aged first?
- ★ Cache-hit speculated last
- ★ Branch speculated first?
- ★ Branches first?




Issue 3: Cache?

- ★ Shared cache?
- ★ Coherence?
- ★ Trashing?




Issue 4: Synchronization?

- ★ lock?
- ★ memory based(atomic instruction)
- ★ acquire (lock)
- ★ release (lock)




Issue 5: Compiler?

- ★ Minimize cache interference
- ★ Latency
- ★ Sharing optimization




When to use SMT?

- ★ Biggest application with shared resources
 - e.g. web server
- ★ What else?



How to multithreading?





Good old POSIX thread (pthread)

```
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include <pthread.h>
#include <unistd.h>

#define NUM_THREADS 5

void *perform_work(void *arguments){
    int index = *((int *)arguments);
    int sleep_time = 1 + rand() % NUM_THREADS;
    printf("THREAD %d: Started.\n", index);
    printf("THREAD %d: Will be sleeping for %d seconds.\n", index, sleep_time);
    sleep(sleep_time);
    printf("THREAD %d: Ended.\n", index);
}
```


```
int main(void) {
    pthread_t threads[NUM_THREADS];
    int thread_args[NUM_THREADS];
    int i;
    int result_code;

    //create all threads one by one
    for (i = 0; i < NUM_THREADS; i++) {
        printf("IN MAIN: Creating thread %d.\n", i);
        thread_args[i] = i;
        result_code = pthread_create(&threads[i], NULL, perform_work, &thread_args[i]);
        assert(!result_code);
    }

    printf("IN MAIN: All threads are created.\n");

    //wait for each thread to complete
    for (i = 0; i < NUM_THREADS; i++) {
        result_code = pthread_join(threads[i], NULL);
        assert(!result_code);
        printf("IN MAIN: Thread %d has ended.\n", i);
    }

    printf("MAIN program has ended.\n");
    return 0;
}
```



OpenMP

```
#pragma omp parallel num_threads(10)
```


- ★ API for Threads, use directive (`#pragma omp`) to hint compiler
- ★ Work with C and Fortran

```
// gcc -fopenmp -o loop loop.c
#include <stdio.h>
#include <omp.h>

int main(void)
{
    #pragma omp parallel for
    for (int i=0;i<100000;i++) {
        a[i]=2*i+i;
        printf("a[%d],%d\n",i, a[i]);
    }
    return 0;
}
```

```
// gcc -fopenmp -o hello hello.c
#include <stdio.h>
#include <omp.h>

int main(void)
{
    #pragma omp parallel
    printf("Hello, world.\n");
    return 0;
}
```



Grand Central Dispatch


Ported to FreeBSD

- ★ Proposed by Apple
- ★ Used since Mac OS X 10.6 and iOS 4
- ★ An implementation of task parallelism based on the thread pool pattern

```
for (i = 0; i < count; i++) {  
    results[i] = do_work(data, i);  
}  
total = summarize(results, count);
```


Apache GCD MPM HTTP is 0.3-0.5 LOCs comparing to other MPMs.

```
dispatch_apply(count, dispatch_get_global_queue(0, 0), ^(size_t i){  
    results[i] = do_work(data, i);  
});  
total = summarize(results, count);
```




Other tools

- ★ Intel Threading Building Blocks
 - C++ template library developed by Intel for Parallel programming on multi-core processors.
- ★ Java Concurrency
- ★ ... etc ...



Exercise






Exercises

- ★ What is the main benefit of openmp over standard thread (eg. POSIX thread)?
- ★ For a given code, modify it to take the benefit of simultaneous multithreading processor using openmp. With the new code, what is the potential speed up?


```
#include <stdio.h>
int main(void)
{
    int a[1000000];
    for (int i=0;i<1000000;i++) {
        a[i]=2*i+i;
        printf("a[%d],%d\n",i, a[i]);
    }
    return 0;
}
```

Exercises

- ★ Base on OpenMP, explain the concepts of work sharing constructs for
 - loop constructs: for and do
 - sections
 - single
 - Workshare

HINT: <https://www.openmp.org/wp-content/uploads/openmp-examples-4.5.0.pdf>



End of Lecture 5

